



Nothing is larger than life.™

# Tech talk May 8th 2019

## Datamon

Datamon manages infinite reflections of data

# Background

- Building a ML/AI app typically includes multiple stages that process data where the stage (compute) and the data have their own life cycles.
- The process to reverse engineering the output variations over time requires insight into every component and how they change over time.
- Insight into variations in components and how it impacts the end result are needed to automate promotion of individual components into production at scale.

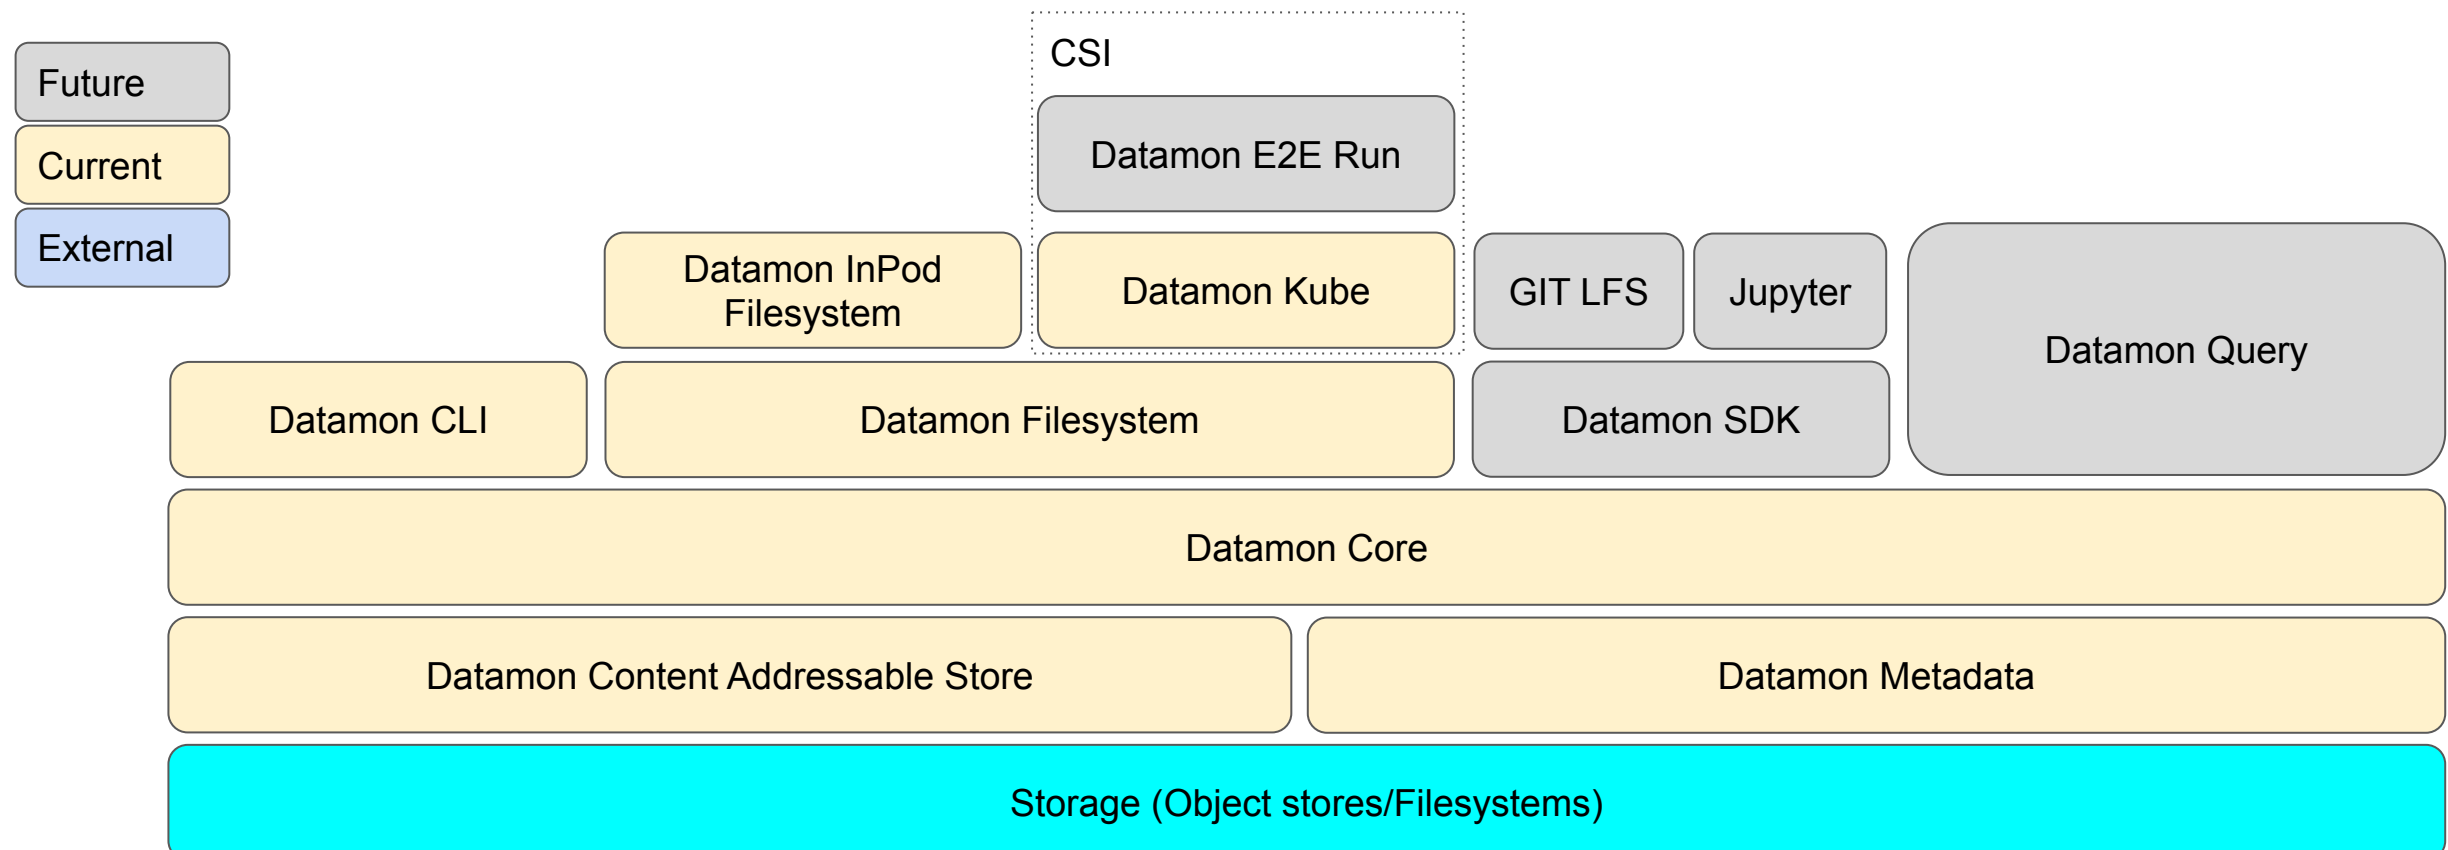
Datamon



# Datamon

Datamon: Git for data + more. Written entirely in GO

- Stores and versions all data under its management
- Understands how a ML pipeline is built and tracks versions of every component in the pipeline (compute + data)
- Can answer the question: What changed?
- Inexpensive and cheap
  - Inexpensive and easy to fork data and run experiments.
    - Deduplication of data
  - 10x cheaper than current solution
- Geo redundancy of data
  - Serverless functionality, geo replicated storage via GCS
- High concurrency of usage with “infinite” bandwidth.

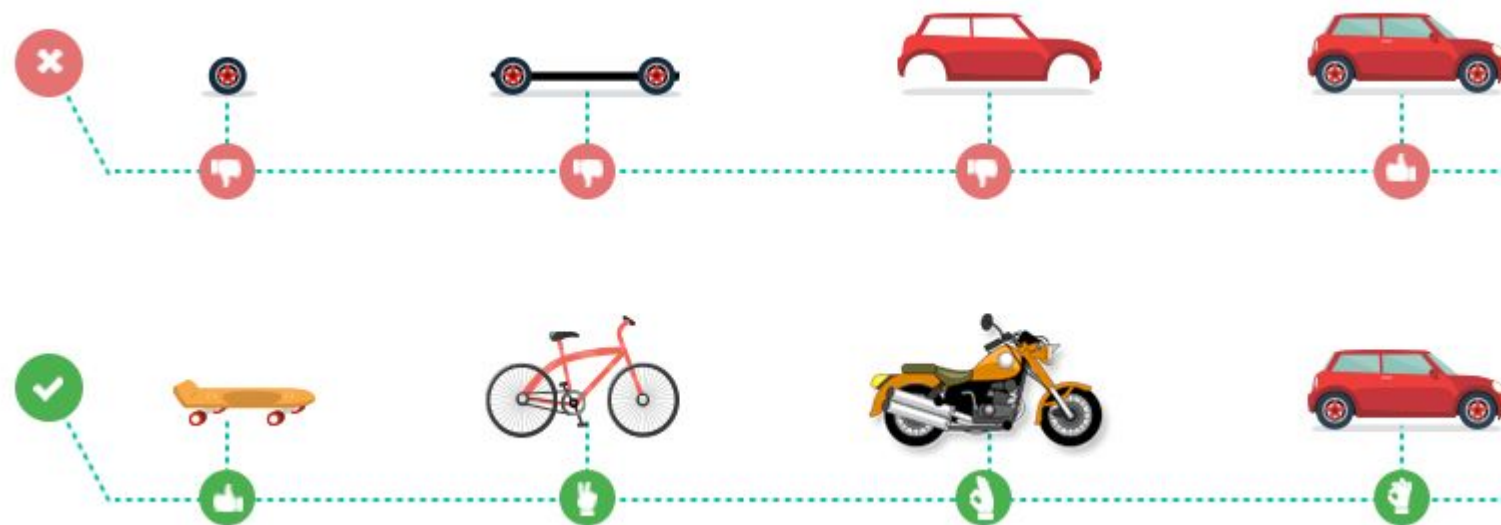


# What does it mean?

- Treat it like git + dropbox
  - Any dataset that needs to be stored, create a personal repo and store it.
- Any python/fortran/go code that needs to process data in K8S use the sample yaml as template for fetching and storing data in Datamon
- On boarding
  - Try the CLI for personal use
  - Over the next few weeks existing workloads deployed in kubernetes will be rejuggled to fetch and store data in Datamon
- It is open source, tell your friends!

# Development model

- Iterative
- Feedback is important
  - Feedback during on boarding will be used to guide the next set of features to work on. So provide feedback!
- You are the customer!
  - File issues!
  - Request features!



Source: <https://www.quicksrum.com/Article/ArticleDetails/5174/1/Why-start-up-should-focus-on-Minimum-Viable-Product>

# Datamon Model

Repo: Repo is a unit of data that needs to exist together.

Bundle: Bundle analogous to git commit is a point in time copy/version of the data in a repo

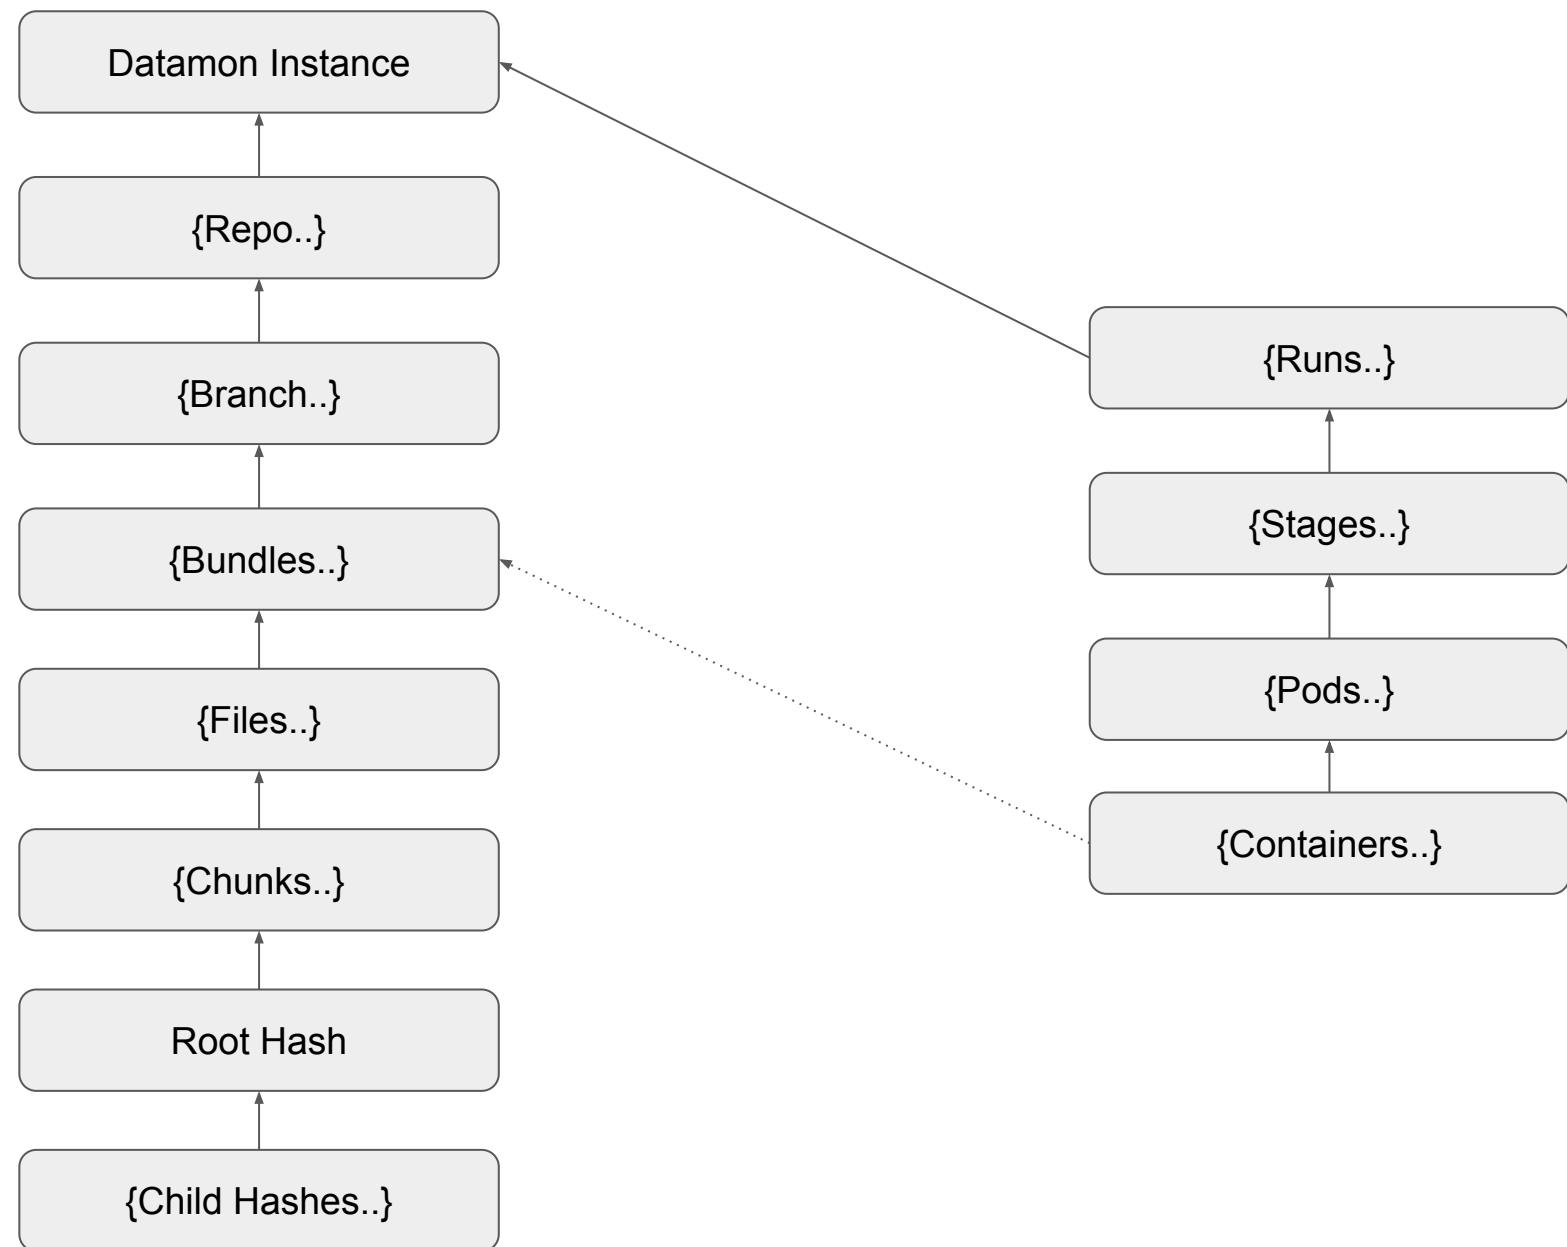
Planned features:

Tags: A name given to a bundle. Example: Latest, production

Branches: Different branches for the data in a repo.

Runs: ML pipeline run metadata that includes the versions of compute and data in use for a given run of a pipeline.

# Datamon Core & Content Addressable Store



{..} Set of





# Datamon CLI

Create repo analogous to git repo

```
datamon repo create --description "Ritesh's repo for testing" --repo ritesh-datamon-test-repo
```

Upload a bundle, the last line prints the commit hash. This will be needed for downloading the bundle

```
#datamon bundle upload --path /path/to/data/folder --message "The initial commit for the repo" --repo ritesh-datamon-test-repo
Uploaded bundle id:1INzQ5TV4vAAfU2PbRFgPfnzEwR
```

List bundles in a repo

```
#datamon bundle list --repo ritesh-test-repo
Using config file: /Users/ritesh/.datamon/datamon.yaml
1INzQ5TV4vAAfU2PbRFgPfnzEwR , 2019-03-12 22:10:24.159704 -0700 PDT , Updating test bundle
```

Download a bundle

```
datamon bundle download --repo ritesh-test-repo --destination /path/to/folder/to/download --bundle 1INzQ5TV4vAAfU2PbRFgPfnzEwR
```

List all files in a bundle

```
datamon bundle list files --repo ritesh-test-repo --bundle 1ISwIzeAR6m3a0VltAsj1kfQaml
```

Download a single file from a bundle

```
datamon bundle download file --file datamon/cmd/repo_list.go --repo ritesh-test-repo --bundle 1ISwIzeAR6m3a0VltAsj1kfQaml
```

# Datamon Filesystem

```
datamon git:(csi) x datamon bundle mount --bundle 1KYmD7EKH7EtovcUn0wLIX9LeqG --destination /tmp/data --mount /tmp/mount --repo ritesh-test-repo
```

```
→ datamon git:(csi) x cd /tmp/mount
```

```
→ mount ls
```

```
backup2blobs csi
```

```
datamon
```

```
→ mount ls -laR
```

```
total 32
```

drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	.
drwxrwxrwt	9	root	wheel	288	May	6	15:21	..
drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	backup2blobs
drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	csi
drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	datamon

```
./backup2blobs:
```

```
total 32
```

drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	.
drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	..
drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	cmd
-rw-r-xr-x	1	root	wheel	272	Apr	29	15:26	main.go

```
./backup2blobs/cmd:
```

```
total 72
```

drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	.
drwxr-xr-x	2	root	wheel	2048	Apr	29	15:26	..
-rw-r-xr-x	1	root	wheel	5657	Apr	29	15:26	blob2file.go
-rw-r-xr-x	1	root	wheel	6296	Apr	29	15:26	file2blobs.go
-rw-r-xr-x	1	root	wheel	8563	Apr	29	15:26	generateFileList.go

# Kubernetes integration

1. Define a storage class that refers to the Datamon Repo

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: datamon
provisioner: com.datamon.csi
parameters:
  repo: "ritesh-test-repo"
```

2. Define a claim that refers to the bundle

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: thankyoumsau42
spec:
  accessModes:
    - ReadOnlyMany
  storageClassName: datamon
  resources:
    requests:
      storage: 1Ti
  selector:
    matchLabels:
      release: "1KYmD7EKH7EtovcUn0wIIX9LeqG"
```

3. Define a pod that refers to the claim

```
apiVersion: v1
kind: Pod
metadata:
  # This name uniquely identifies the Deployment
  name: datamon-dev
spec:
  containers:
    - name: dev
      image: golang:1.11-alpine
      tty: true
      stdin: true
      volumeMounts:
        - name: credentials
          readOnly: true
          mountPath: "/etc/datamon-creds"
        - name: riteshpv2
          readOnly: true
          mountPath: "/data"
  volumes:
    - name: credentials
      secret:
        secretName: gcs-credentials
    - name: riteshpv2
      persistentVolumeClaim:
        claimName: thankyoumsau42
        readOnly: true
```



# Kubernetes integration

```
→ mount kubectl attach --namespace=dev -it datamon-dev -c dev
If you don't see a command prompt, try pressing enter.
/go #
/go #
/go #
/go # cd /data/
/data # ls
backup2blobs  csi          datamon
/data # ls -laR
.:
total 6
drwxr-xr-x    2 root    root          2048 Apr 29 22:26 backup2blobs
drwxr-xr-x    2 root    root          2048 Apr 29 22:26 csi
drwxr-xr-x    2 root    root          2048 Apr 29 22:26 datamon

./backup2blobs:
total 3
drwxr-xr-x    2 root    root          2048 Apr 29 22:26 cmd
-rw-r-xr-x    1 root    root           272 Apr 29 22:26 main.go

./backup2blobs/cmd:
total 21
-rw-r-xr-x    1 root    root          5657 Apr 29 22:26 blob2file.go
-rw-r-xr-x    1 root    root          6296 Apr 29 22:26 file2blobs.go
-rw-r-xr-x    1 root    root          8563 Apr 29 22:26 generateFileList.go

./csi:
total 1
-rw-r-xr-x    1 root    root           131 Apr 29 22:26 main.go

./datamon:
total 3
drwxr-xr-x    2 root    root          2048 Apr 29 22:26 cmd
-rw-r-xr-x    1 root    root           301 Apr 29 22:26 main.go
```

# Datamon InPod Filesystem

1. Add a volume to contain mountpoints

```
volumes:  
- name: fuse-mountpoint  
  emptyDir: {}  
- name: google-application-credentials  
  secret:  
    secretName: google-application-credentials
```

2. Add a InPod Filesystem container entry that provides bundles to the volume

```
- name: datamon-sidecar  
  image: gcr.io/onec-co/datamon-fuse-demo-sidecar:latest  
  imagePullPolicy: "Always"  
  command: ["datamon"]  
  args: ["bundle", "mount", "--repo", "ransom-datamon-test-repo", "--destination", "/tmp", "--mount", "/tmp/mount"]  
  securityContext:  
    privileged: true  
  stdin: true  
  tty: true  
  volumeMounts:  
  - mountPath: /tmp/mount  
    name: fuse-mountpoint  
    mountPropagation: "Bidirectional"  
  - mountPath: /tmp/gac  
    name: google-application-credentials  
  env:  
  - name: GOOGLE_APPLICATION_CREDENTIALS  
    value: /tmp/gac/google-application-credentials.json
```

3. Add a volume entry to the application container

```
volumeMounts:  
- mountPath: /tmp/mount  
  name: fuse-mountpoint  
  mountPropagation: "HostToContainer"
```

# Datamon InPod Filesystem

Then the bundle is available via filesystem operations in the application container

```
datamon% kubectl create -f hack/k8s/gen/example-ro.yaml
deployment.apps/datamon-ro-demo created
datamon% kubectl exec -it datamon-ro-demo-689b69bf84-s42xm -c demo-shell -- "/bin/bash"
root@datamon-ro-demo-689b69bf84-s42xm:/# ls /tmp/mount
create_ro_pod.sh  datamon.yaml  run_shell.sh  shell.Dockerfile  sidecar.Dockerfile
root@datamon-ro-demo-689b69bf84-s42xm:/#
```

# Demo

1. Kubernetes Volume
2. CLI



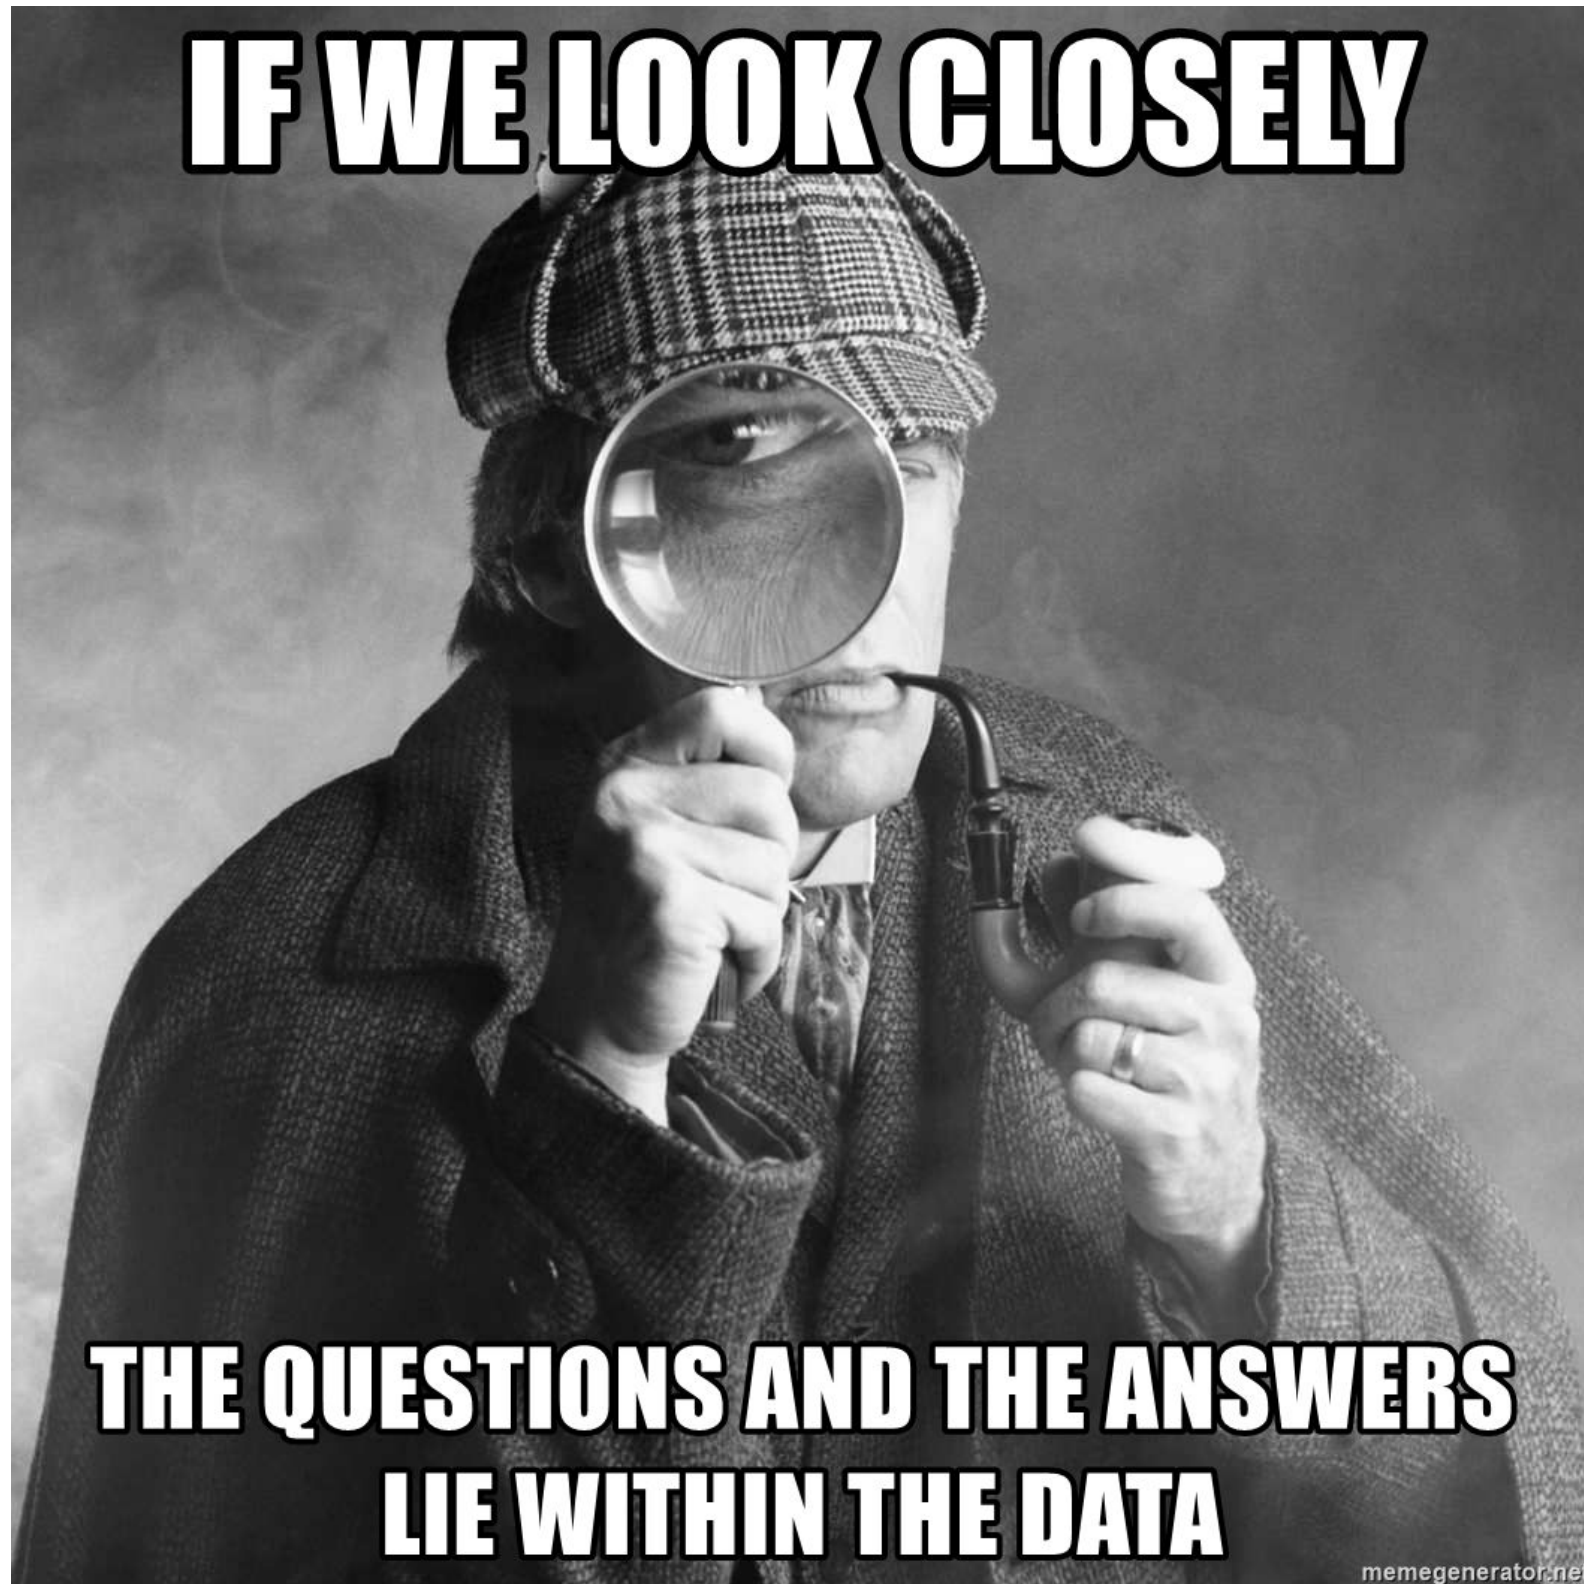
# Next Steps: Integration & Query

1. Existing pipelines moved to datamon
  - a. Integrate
  - b. Measure
  - c. Improve
2. Track metadata for runs
  - a. The metadata for repo/bundle/branch is linked for a pipeline
3. Track quality of output and allow experimentations at scale.
  - a. Metadata allows for queries that span the pipeline to allow the following questions to be answered
    - i. What changed between 2 runs?
    - ii. If only one component is changed, how does the result vary?
4. Metadata tracking and query is a key value add beyond version management, efficiency and cost savings

# What is next?

- New release coming next week
  - Kubernetes
  - InPod FS
  - More testing and bug fixes.
- Migration of pipelines:
  - Integration with existing pipelines
- Performance measurements and performance improvements
  - Where are the pain points that matter to OneConcern
- Argo integration
  - End to end visibility
- Datamon query
  - Query end to end visibility
  - Dive deep into changes over time
- Access control and encryption

**IF WE LOOK CLOSELY**



**THE QUESTIONS AND THE ANSWERS  
LIE WITHIN THE DATA**

# Similar Projects?

- Pachyderm
- DVC.org
  - Deduplication of data at rest
  - Streaming and caching of data
  - Kubernetes integration
  - Tightly coupled with GIT